

DATA200 Final Project

Preliminary Work

```
# Load necessary libraries
library(tidyverse)
```

```
## — Attaching core tidyverse packages — tidyverse 2.0.0 —
## ✓ dplyr      1.1.4      ✓ readr      2.1.5
## ✓ forcats    1.0.0      ✓ stringr    1.5.1
## ✓ ggplot2    3.5.1      ✓ tibble     3.2.1
## ✓ lubridate  1.9.3      ✓ tidyr      1.3.1
## ✓ purrr      1.0.2
## — Conflicts — tidyverse_conflicts() —
## ✖ dplyr::filter() masks stats::filter()
## ✖ dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts
to become errors
```

```
library(ggplot2)
library(dplyr)
library(caret)
```

```
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##     lift
```

```
# Set working directory -> change path url as needed!!!
# *if this isn't working, manually import data.
setwd('/Users/hannahmarr/Desktop/Tufts/DATA200/Final_Project_200/')

# Import data
acceptances <- read.csv("DisasterDeclarationsSummaries.csv")
denials <- read.csv("DeclarationDenials.csv")
```

Exploratory Analysis

```
#### FOR 'acceptances' DATASET ####
# Display the first few rows of the dataframe to inspect the data
head(acceptances)
```

femaDeclarationString <chr>	disasterNumber <int>	state <chr>	declarationType <chr>	declarationDate <chr>
1 FM-5530-NV	5530	NV	FM	2024-08-12T00:00:00.000Z
2 FM-5529-OR	5529	OR	FM	2024-08-09T00:00:00.000Z
3 FM-5528-OR	5528	OR	FM	2024-08-06T00:00:00.000Z
4 FM-5527-OR	5527	OR	FM	2024-08-02T00:00:00.000Z
5 FM-5526-CO	5526	CO	FM	2024-08-01T00:00:00.000Z
6 FM-5525-CO	5525	CO	FM	2024-07-31T00:00:00.000Z

6 rows | 1-7 of 29 columns

```
# Get the dimensions of the dataframe (number of rows and columns).
dim(acceptances)
```

```
## [1] 67245    28
```

```
# Check for missing values in each column.
colSums(is.na(acceptances))
```

```
##      femaDeclarationString      disasterNumber      state
##              0              0              0
##      declarationType      declarationDate      fyDeclared
##              0              0              0
##      incidentType      declarationTitle      ihProgramDeclared
##              0              0              0
##      iaProgramDeclared      paProgramDeclared      hmProgramDeclared
##              0              0              0
##      incidentBeginDate      incidentEndDate      disasterCloseoutDate
##              0              0              0
##      tribalRequest      fipsStateCode      fipsCountyCode
##              0              0              0
##      placeCode      designatedArea      declarationRequestNumber
##              0              0              0
##      lastIAFilingDate      incidentId      region
##              0              0              0
##      designatedIncidentTypes      lastRefresh      hash
##              0              0              0
##              id
##              0
```

```
#### FOR 'denials' DATASET #####
# Display the first few rows of the dataframe to inspect the data
head(denials)
```

	declarationRequestNumber <int>	region <int>	stateAbbreviation <chr>	state <chr>	tribalRequest <int>
1	23	10	ID	Idaho	0
2	40	6	OK	Oklahoma	0
3	46	6	LA	Louisiana	0
4	51	4	NC	North Carolina	0
5	56	5	IL	Illinois	0
6	79	6	TX	Texas	0

6 rows | 1-6 of 21 columns

```
# Get the dimensions of the dataframe (number of rows and columns).
dim(denials)
```

```
## [1] 317 20
```

```
# Check for missing values in each column.
colSums(is.na(denials))
```

```
## declarationRequestNumber      region
##                0                0
##      stateAbbreviation      state
##                0                0
##      tribalRequest      declarationRequestDate
##                0                0
##      declarationRequestType      incidentName
##                0                0
##      requestedIncidentTypes      requestedIncidentBeginDate
##                0                0
##      requestedIncidentEndDate      currentRequestStatus
##                0                0
##      requestStatusDate      ihProgramRequested
##                0                0
##      iaProgramRequested      paProgramRequested
##                0                0
##      hmProgramRequested      incidentId
##                0                0
##      incidentBeginDate      id
##                0                0
```

We will not immediately drop rows with null values, since values in other columns in these rows may become applicable, while the columns with null values do not need to be included in analysis.

Cleaning the Data

First, we did a manual inspection of the column names of the 'denials' dataset and adjusted them so that the two datasets could be merged

```
denials <- select(denials, -state) #orig, 'state' was full name not abbreviation
```

```
denials <- rename(denials,
                  state = stateAbbreviation,
                  declarationDate = declarationRequestDate,
                  declarationType = declarationRequestType,
                  declarationTitle = incidentName,
                  incidentType = requestedIncidentTypes,
                  incidentEndDate = requestedIncidentEndDate,
                  )
```

```
str(denials)
```

```
## 'data.frame':   317 obs. of  19 variables:
## $ declarationRequestNumber : int  23 40 46 51 56 79 93 138 154 155 ...
## $ region                  : int  10 6 6 4 5 6 1 6 6 5 ...
## $ state                   : chr  "ID" "OK" "LA" "NC" ...
## $ tribalRequest           : int  0 0 0 0 0 0 0 0 0 0 ...
## $ declarationDate         : chr  "2000-03-03T00:00:00.000Z" "2000-05-10T00:00:00.0
00Z" "2000-06-05T00:00:00.000Z" "2000-06-06T00:00:00.000Z" ...
## $ declarationType         : chr  "Major Disaster" "Major Disaster" "Major Disaste
r" "Major Disaster" ...
## $ declarationTitle        : chr  "ID-ROCKSLIDES-02-11-2000" "OK-Tulsa Flooding-5-6
-00" "LA-Tornado-North Louisiana-4/23/00" "NC -Severe Storms and Possible Tornadoes" ...
## $ incidentType            : chr  "Mud/Landslide" "Flood" "Tornado" "Severe Storm"
...
## $ requestedIncidentBeginDate: chr  "2000-01-30T00:00:00.000Z" "2000-05-05T00:00:00.0
00Z" "2000-04-23T00:00:00.000Z" "2000-05-25T00:00:00.000Z" ...
## $ incidentEndDate         : chr  "" "2000-05-08T00:00:00.000Z" "2000-04-24T00:00:0
0.000Z" "2000-05-25T00:00:00.000Z" ...
## $ currentRequestStatus    : chr  "Turndown" "Turndown" "Turndown" "Turndown" ...
## $ requestStatusDate       : chr  "2000-03-17T00:00:00.000Z" "2000-05-24T00:00:00.0
00Z" "2000-06-21T00:00:00.000Z" "2000-06-23T00:00:00.000Z" ...
## $ ihProgramRequested      : int  0 0 0 0 0 0 0 0 0 0 ...
## $ iaProgramRequested      : int  1 1 1 0 0 0 1 0 0 0 ...
## $ paProgramRequested      : int  0 0 0 1 1 1 0 1 1 1 ...
## $ hmProgramRequested      : int  1 1 1 1 1 0 1 0 0 1 ...
## $ incidentId              : int  2000021103 2000050701 2000042401 2000052601 20000
51901 2000073101 2000081404 2000090803 2000092205 2000091401 ...
## $ incidentBeginDate       : chr  "2000-01-30T00:00:00.000Z" "2000-05-06T00:00:00.0
00Z" "2000-04-23T00:00:00.000Z" "2000-05-25T00:00:00.000Z" ...
## $ id                      : chr  "61ebab68-adda-44c9-bca0-7c3b7a87c40f" "cf5fd28e-
0195-43be-9bbf-9f153dcc5b40" "779546f1-d286-45de-95f3-ff8905723a9f" "9c6e0548-dd06-48f5-
9d7a-32dc6ccc24d5" ...
```

```
str(acceptances)
```

```
## 'data.frame':    67245 obs. of  28 variables:
## $ femaDeclarationString : chr "FM-5530-NV" "FM-5529-OR" "FM-5528-OR" "FM-5527-OR"
...
## $ disasterNumber       : int  5530 5529 5528 5527 5526 5525 5525 5524 5523 5522
...
## $ state                : chr  "NV" "OR" "OR" "OR" ...
## $ declarationType      : chr  "FM" "FM" "FM" "FM" ...
## $ declarationDate      : chr  "2024-08-12T00:00:00.000Z" "2024-08-09T00:00:00.000Z" "2024-08-06T00:00:00.000Z" "2024-08-02T00:00:00.000Z" ...
## $ fyDeclared          : int   2024 2024 2024 2024 2024 2024 2024 2024 2024 2024
...
## $ incidentType        : chr  "Fire" "Fire" "Fire" "Fire" ...
## $ declarationTitle     : chr  "GOLD RANCH FIRE" "LEE FALLS FIRE" "ELK LANE FIRE" "MILE MARKER 132 FIRE" ...
## $ ihProgramDeclared   : int    0 0 0 0 0 0 0 0 0 0 ...
## $ iaProgramDeclared   : int    0 0 0 0 0 0 0 0 0 0 ...
## $ paProgramDeclared   : int    1 1 1 1 1 1 1 1 1 1 ...
## $ hmProgramDeclared   : int    1 1 1 1 1 1 1 1 1 1 ...
## $ incidentBeginDate   : chr  "2024-08-11T00:00:00.000Z" "2024-08-08T00:00:00.000Z" "2024-08-04T00:00:00.000Z" "2024-08-02T00:00:00.000Z" ...
## $ incidentEndDate     : chr  "" "" "" "" ...
## $ disasterCloseoutDate : chr  "" "" "" "" ...
## $ tribalRequest       : int    0 0 0 0 0 0 0 0 0 0 ...
## $ fipsStateCode       : int   32 41 41 41 8 8 8 8 56 6 ...
## $ fipsCountyCode      : int   31 67 31 17 59 13 69 69 31 29 ...
## $ placeCode           : int  99031 99067 99031 99017 99059 99013 99069 99069 99031 99029 ...
## $ designatedArea      : chr  "Washoe (County)" "Washington (County)" "Jefferson (County)" "Deschutes (County)" ...
## $ declarationRequestNumber: int  24123 24122 24116 24111 24106 24105 24105 24104 24103 24102 ...
## $ lastIAFilingDate     : chr  "" "" "" "" ...
## $ incidentId           : num  2.02e+09 2.02e+09 2.02e+09 2.02e+09 2.02e+09 ...
## $ region               : int   9 10 10 10 8 8 8 8 8 9 ...
## $ designatedIncidentTypes : chr  "R" "R" "R" "R" ...
## $ lastRefresh          : chr  "2024-08-27T18:22:14.800Z" "2024-08-27T18:22:14.800Z" "2024-08-27T18:22:14.800Z" "2024-08-27T18:22:14.800Z" ...
## $ hash                 : chr  "5d07e7c51bb300bfbec94a699a1e1ab1d61a97cd" "ae87cf3c6ed795015b714af7166c7c295b2b67c7" "432cf0995c47e3895cea696ede5621b810460501" "2f21d90cb6bc64b0d4121aa3f18d852bbb4b11fa" ...
## $ id                   : chr  "f15a7a79-f1c3-41bb-8a5c-c05fbae34423" "09e3f81a-5e16-4b72-b317-1c64e0cfa59c" "59983f89-30bf-4888-b21b-62e8d57d9aac" "8d13ecf0-bc2f-496b-8c9f-b2e73da832a0" ...
```

Merging the Datasets

```
# Merging the two datasets
data <- merge(denials, acceptances, by = c("state",
                                           "tribalRequest",
                                           "declarationDate",
                                           "declarationTitle",
                                           "incidentType",
                                           "incidentEndDate",
                                           "incidentId",
                                           "incidentBeginDate",
                                           "id",
                                           "region",
                                           "declarationType"), all = TRUE)
```

```
# Inspect the dataset to determine if the merge was successful
head(data)
```

state <chr>	tribalRequest <int>	declarationDate <chr>	declarationTitle <chr>	incidentType <chr>
1 AK	0	1953-10-30T00:00:00.000Z	SEVERE HARDSHIP	Other
2 AK	0	1954-11-10T00:00:00.000Z	SEVERE HARDSHIP	Other
3 AK	0	1955-12-22T00:00:00.000Z	SEVERE HARDSHIP	Other
4 AK	0	1964-03-28T00:00:00.000Z	EARTHQUAKE	Earthquake
5 AK	0	1967-08-17T00:00:00.000Z	SEVERE STORMS & FLOODING	Flood
6 AK	0	1969-12-19T00:00:00.000Z	HEAVY RAINS & LANDSLIDE	Severe Storm

6 rows | 1-6 of 37 columns

```
dim(data)
```

```
## [1] 67562    36
```

```
str(data)
```

```
## 'data.frame':    67562 obs. of  36 variables:
## $ state                : chr  "AK" "AK" "AK" "AK" ...
## $ tribalRequest        : int   0 0 0 0 0 0 0 0 0 0 ...
## $ declarationDate      : chr  "1953-10-30T00:00:00.000Z" "1954-11-10T00:00:00.000Z" "1955-12-22T00:00:00.000Z" "1964-03-28T00:00:00.000Z" ...
## $ declarationTitle     : chr  "SEVERE HARDSHIP" "SEVERE HARDSHIP" "SEVERE HARDSHIP" "EARTHQUAKE" ...
## $ incidentType         : chr  "Other" "Other" "Other" "Earthquake" ...
## $ incidentEndDate      : chr  "1953-10-30T00:00:00.000Z" "1954-11-10T00:00:00.000Z" "1955-12-22T00:00:00.000Z" "1964-03-28T00:00:00.000Z" ...
## $ incidentId           : num   53012 54019 55017 64015 67024 ...
## $ incidentBeginDate    : chr  "1953-10-30T00:00:00.000Z" "1954-11-10T00:00:00.000Z" "1955-12-22T00:00:00.000Z" "1964-03-28T00:00:00.000Z" ...
## $ id                   : chr  "353cdf38-9ad2-40cd-b92a-58976c13236f" "6cf23fc4-1cd6-4d3d-8dd5-154ab481d09f" "8a2aae98-dadd-4b5d-b449-ae238dba5d3" "6fba469d-7108-4c1e-b82d-60548ede4178" ...
## $ region               : int   10 10 10 10 10 10 10 10 10 10 ...
## $ declarationType      : chr  "DR" "DR" "DR" "DR" ...
## $ declarationRequestNumber.x: int   NA NA NA NA NA NA NA NA NA NA ...
## $ requestedIncidentBeginDate: chr  NA NA NA NA ...
## $ currentRequestStatus  : chr  NA NA NA NA ...
## $ requestStatusDate     : chr  NA NA NA NA ...
## $ ihProgramRequested    : int   NA NA NA NA NA NA NA NA NA NA ...
## $ iaProgramRequested    : int   NA NA NA NA NA NA NA NA NA NA ...
## $ paProgramRequested    : int   NA NA NA NA NA NA NA NA NA NA ...
## $ hmProgramRequested    : int   NA NA NA NA NA NA NA NA NA NA ...
## $ femaDeclarationString : chr  "DR-13-AK" "DR-31-AK" "DR-46-AK" "DR-168-AK" ...
## $ disasterNumber       : int   13 31 46 168 230 281 2001 2004 2005 2006 ...
## $ fyDeclared           : int   1954 1955 1956 1964 1967 1970 1970 1971 1971 1973 ...
## $ ihProgramDeclared    : int   0 0 0 0 0 0 0 0 0 0 ...
## $ iaProgramDeclared    : int   1 1 1 1 1 1 0 0 0 0 ...
## $ paProgramDeclared    : int   1 1 1 1 1 1 1 1 1 1 ...
## $ hmProgramDeclared    : int   1 1 1 1 0 0 0 0 0 0 ...
## $ disasterCloseoutDate  : chr  "1957-09-01T00:00:00.000Z" "1957-09-01T00:00:00.000Z" "1956-12-01T00:00:00.000Z" "1971-06-17T00:00:00.000Z" ...
## $ fipsStateCode        : int   2 2 2 2 2 2 2 2 2 2 ...
## $ fipsCountyCode       : int   0 0 0 0 0 0 0 0 0 0 ...
## $ placeCode            : int   0 0 0 0 75165 75165 0 0 0 0 ...
## $ designatedArea       : chr  "Statewide" "Statewide" "Statewide" "Statewide" ...
## $ declarationRequestNumber.y: int   53012 54019 55017 64015 67024 69044 70100 71042 71044 73050 ...
## $ lastIAFilingDate     : chr  "" "" "" "" ...
## $ designatedIncidentTypes : chr  "" "" "" "" ...
## $ lastRefresh          : chr  "2024-08-27T18:22:14.800Z" "2024-08-27T18:22:14.800Z" "2024-08-27T18:22:14.800Z" "2024-08-27T18:22:14.800Z" ...
## $ hash                 : chr  "9c2d55e534115516cf862eed47585d695cd1e24" "8aa70d2832920cf1023d7e14d7a03160f6f191ae" "552c686cac2fba0094e5fc89069778b4a8fd1fd3" "4ab30e317c3672060e315922b972311ec0714ff9" ...
```

More Data Cleaning

We will add a column that designates from which dataset the data came. This column will encode the data as 'Denial' or 'Acceptance', based on if it was in the DeclarationDenials dataset ('denials') or the DisasterDeclarationSummaries dataset ('acceptances').

```
# Create the new column 'requestResult'
data$requestResult <- ifelse(is.na(data$paProgramRequested),
                             'Acceptance',
                             'Denial')
```

Denials data only goes back to 2000, and tribal nations were only allowed to submit their own declaration requests to FEMA in January 2013. Therefore, we will drop all data from 2012 and prior.

```
# Converting the incidentBeginDate column to a date-time format and remove rows with data from before the year 2013.
data$incidentBeginDate <- as.POSIXct(data$incidentBeginDate, format = "%Y-%m-%dT%H:%M:%SZ", tz = "UTC")
data <- data[is.na(data$incidentBeginDate) | format(data$incidentBeginDate, "%Y") >= 2013, ]
```

```
# Inspect data to confirm no issues so far
head(data)
```

state <chr>	tribalRequest <int>	declarationDate <chr>	declarationTitle <chr>	incidentType <chr>
152 AK	0	2013-06-25T00:00:00.000Z	FLOODING	Flood
153 AK	0	2013-06-25T00:00:00.000Z	FLOODING	Flood
154 AK	0	2013-06-25T00:00:00.000Z	FLOODING	Flood
155 AK	0	2013-06-25T00:00:00.000Z	FLOODING	Flood
156 AK	0	2013-06-25T00:00:00.000Z	FLOODING	Flood
157 AK	0	2014-01-16T00:00:00.000Z	FLOODING	Flood

6 rows | 1-6 of 38 columns

```
dim(data)
```

```
## [1] 24439 37
```

Now I will create a new column, declarationMonth, that extracts the month from the declarationDate column to specify the month in which the declaration occurred.

```
data$declarationDate <- as.POSIXct(data$declarationDate, format = "%Y-%m-%dT%H:%M:%SZ",
                                     tz = "UTC")

data$declarationMonth <- format(data$declarationDate, "%m")
```


We will do the same thing for year and create a new column, `declarationYear`, that extracts the year from the `declarationDate` column to specify the year in which the declaration occurred.

```
data$declarationDate <- as.POSIXct(data$declarationDate, format = "%Y-%m-%dT%H:%M:%SZ",
tz = "UTC")

data$declarationYear <- format(data$declarationDate, "%Y")
```

```
# View the first few rows and dimensions of the updated dataset
head(data)
```

state <chr>	tribalRequest <int>	declarationDate <dtm>	declarationTitle <chr>	incidentType <chr>	incidentEndDate <chr>
152 AK	0	2013-06-25	FLOODING	Flood	2013-06-11T00:00:
153 AK	0	2013-06-25	FLOODING	Flood	2013-06-11T00:00:
154 AK	0	2013-06-25	FLOODING	Flood	2013-06-11T00:00:
155 AK	0	2013-06-25	FLOODING	Flood	2013-06-11T00:00:
156 AK	0	2013-06-25	FLOODING	Flood	2013-06-11T00:00:
157 AK	0	2014-01-16	FLOODING	Flood	2013-10-28T00:00:

6 rows | 1-7 of 40 columns

```
dim(data)
```

```
## [1] 24439    39
```

```
# Writing dataset to a csv file to save locally
write.csv(data, "data.csv", row.names = FALSE)
```

Visualizing Acceptances v. Denials for FEMA Aid

```
# Create a summary of the 'requestResult' column
result_summary <- table(data$requestResult)

# Convert the table to a data frame for plotting
result_df <- as.data.frame(result_summary)
colnames(result_df) <- c("RequestResult", "Count")

# Create dynamic legend labels with counts
legend_labels <- paste0(result_df$requestResult, " (", result_df$Count, ")")

# Plot the data using ggplot2
ggplot(result_df, aes(x = RequestResult, y = Count, fill = RequestResult)) +
  geom_bar(stat = "identity") +
  labs(
    title = "Comparison of Acceptances vs. Denials of FEMA aid",
    x = "Request Result",
    y = "Count",
    fill = "Request Result"
  ) +
  scale_fill_discrete(labels = legend_labels) + # Add custom legend labels
  theme_minimal()
```

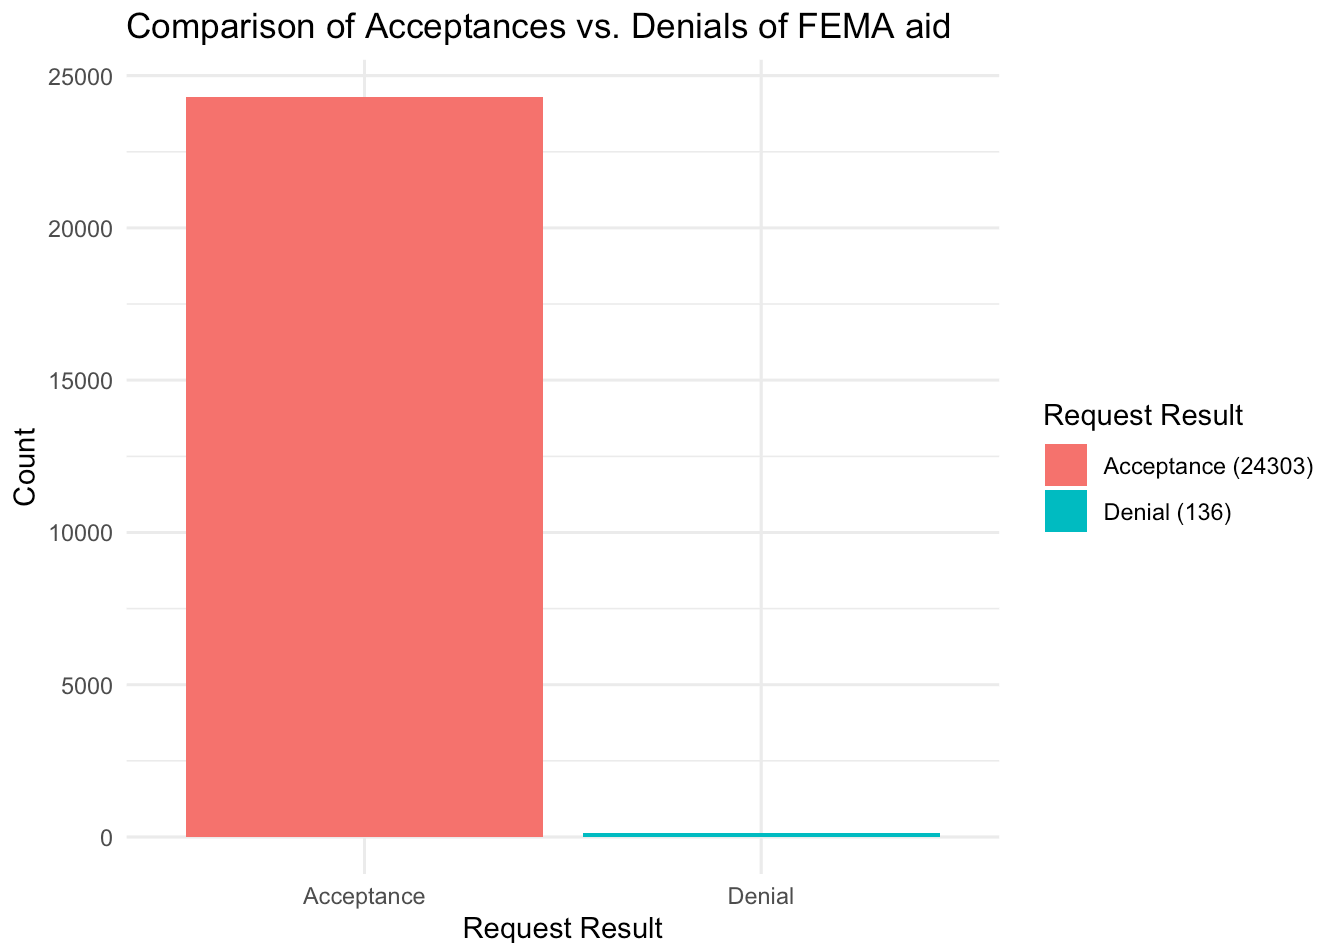


Figure 1. Comparing Overall Acceptances vs. Denials of FEMA aid, we see there are far more documented acceptances than denials of FEMA aid.

Visualizing Variations Among Incident Types

```
# Create a summary of the 'incidentType' column
incident_summary <- table(data$incidentType)

# Convert the summary table to a data frame for visualization
incident_df <- as.data.frame(incident_summary)
colnames(incident_df) <- c("IncidentType", "Count")

# Plot the data using ggplot2
ggplot(incident_df, aes(x = reorder(IncidentType, -Count), y = Count, fill = IncidentType)) +
  geom_bar(stat = "identity") +
  labs(title = "Counts of Each Incident Type", x = "Incident Type", y = "Count") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

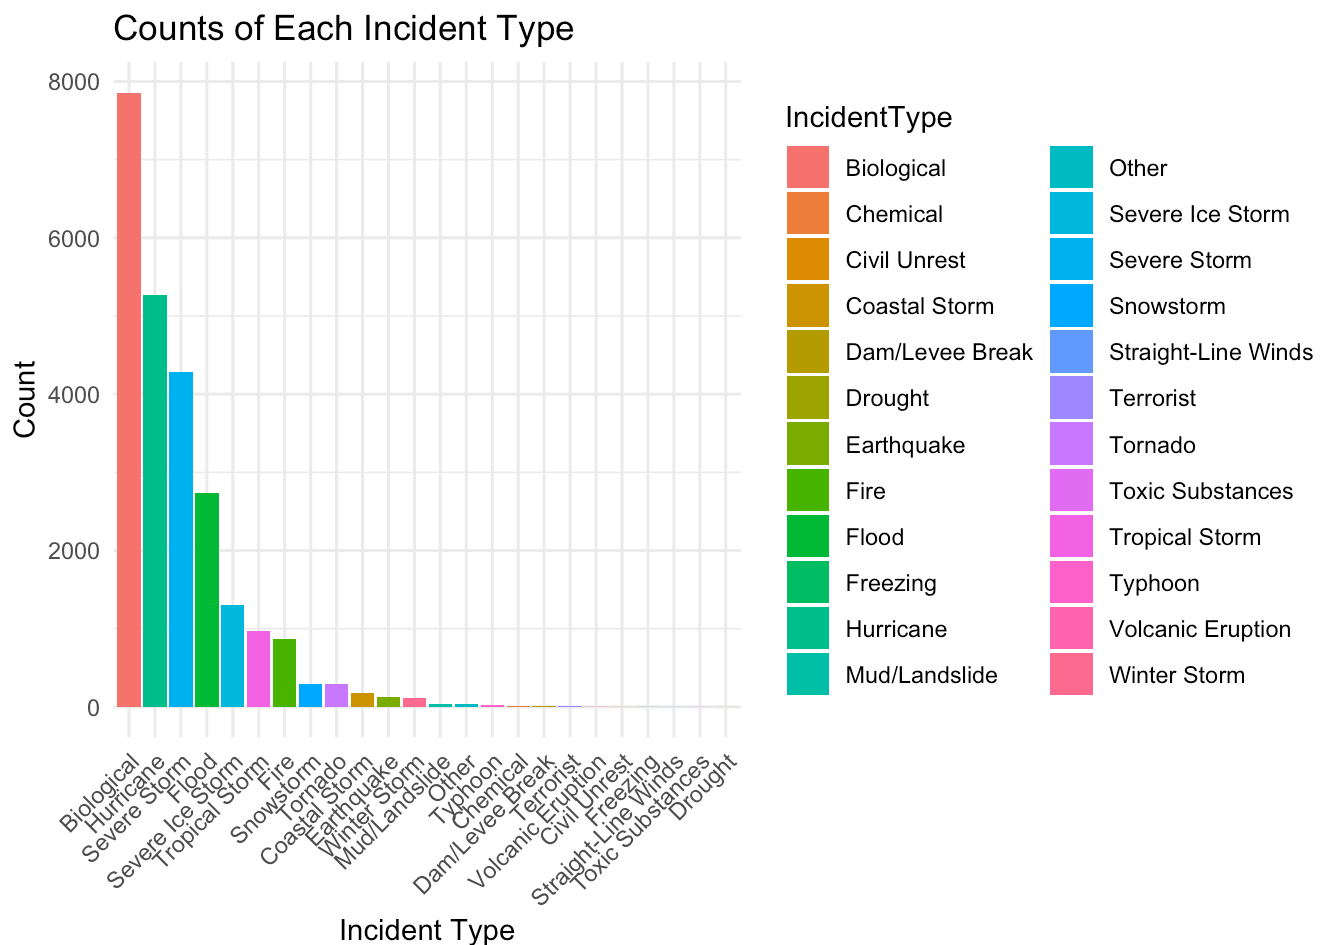


Figure 2. Visualizing the variation among incident types allows us to get a better sense of the various reasons for FEMA aid requests. Biological is the most prevalent incident type, followed by Hurricane, Severe Storm, and Flood.

Examine number of tribal vs. non-tribal nation requests

```

# Create a summary of the tribalRequest column (count 1s and 0s)
tribal_summary <- table(data$tribalRequest)

# Convert the summary table to a data frame for plotting
tribal_df <- as.data.frame(tribal_summary)
colnames(tribal_df) <- c("RequestType", "Count")

# Map 1 to 'Tribal Requests' and 0 to 'Non-Tribal Requests'
tribal_df$RequestType <- factor(tribal_df$RequestType,
                                levels = c(0, 1),
                                labels = c("Non-Tribal Requests", "Tribal Requests"))

# Create custom labels for the legend with counts
custom_labels <- paste(tribal_df$RequestType, " (", tribal_df$Count, ")", sep = "")

# Plot the data using ggplot2
ggplot(tribal_df, aes(x = RequestType, y = Count, fill = RequestType)) +
  geom_bar(stat = "identity") +
  scale_fill_manual(values = c("#4DAF4A", "#377EB8"), labels = custom_labels) +
  labs(title = "Comparison of Tribal vs. Non-Tribal Requests",
       x = "Request Type",
       y = "Count",
       fill = "Request Type (Count)") +
  theme_minimal()

```

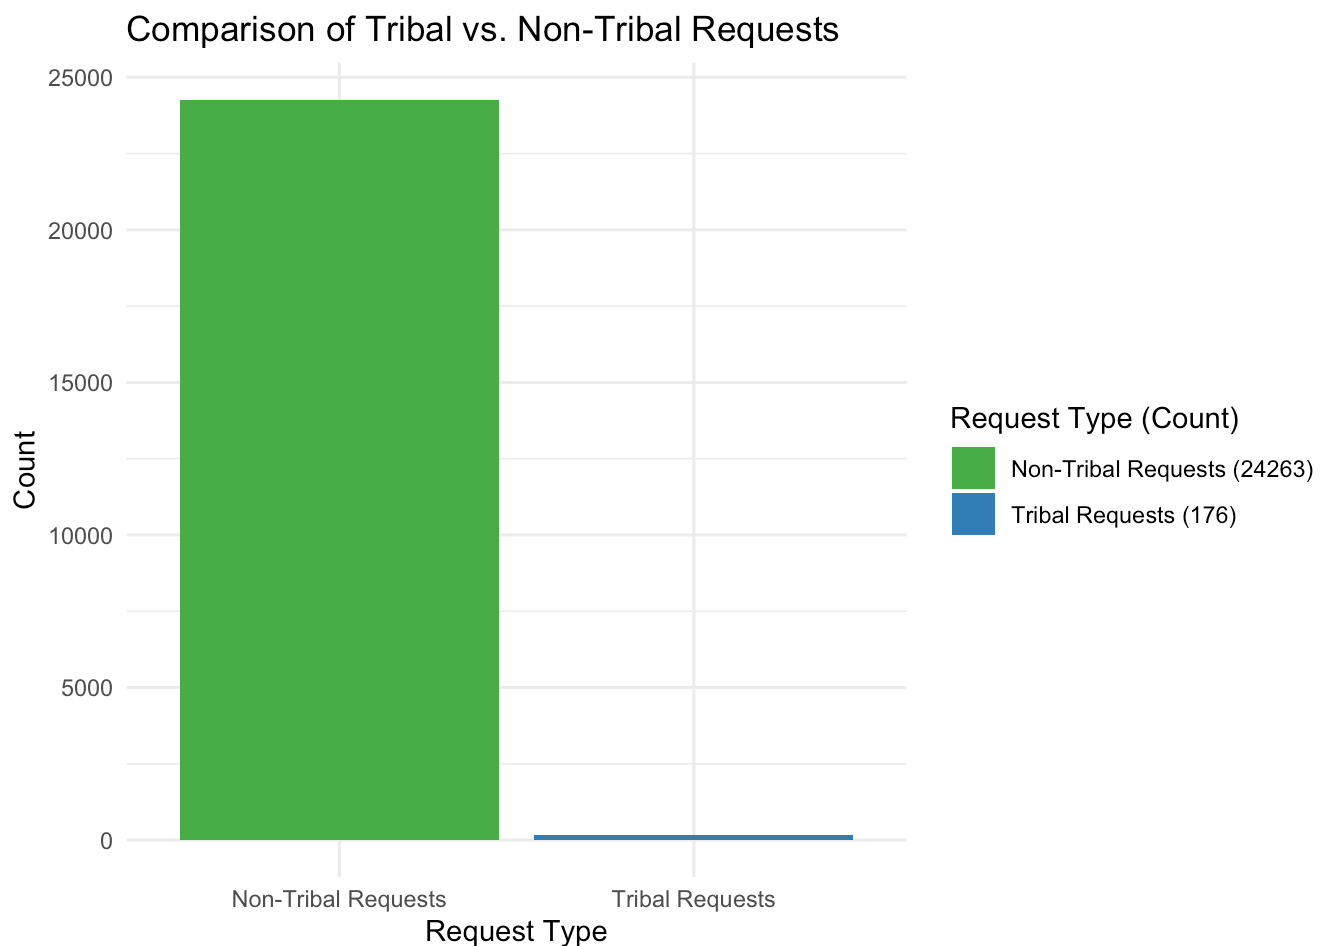


Figure 3. Examining the number of tribal nation vs. non-tribal nation requests shows how infrequent tribal requests are. Part of this divide is likely due to tribal nations not being able to apply for their own declaration request until January 29, 2013. While all pre-2013 data was removed, it is likely that this process is still not as streamlined as the request process for states.

Note that there are very few tribal requests. Later, we will be using LOOCV to ensure results are robust.

Examine number of accepted aid requests for tribal nations over time

```
# Filter data for tribal requests with 'Acceptance' in requestResult
tribal_acceptances <- subset(data, tribalRequest == 1 & requestResult == "Acceptance")

# Count acceptances by year
acceptance_by_year <- table(tribal_acceptances$declarationYear)
acceptance_df <- as.data.frame(acceptance_by_year)
colnames(acceptance_df) <- c("Year", "Count")

# Convert Year to a numeric type for plotting
acceptance_df$Year <- as.numeric(as.character(acceptance_df$Year))

# Plot the data
ggplot(acceptance_df, aes(x = Year, y = Count)) +
  geom_line(color = "blue") +
  geom_point(color = "blue") +
  labs(title = "FEMA Aid Acceptances to Tribal Nations Over Time",
       x = "Year",
       y = "Number of Acceptances") +
  theme_minimal()
```

FEMA Aid Acceptances to Tribal Nations Over Time




Figure 4.1

Examine number of denied aid requests for tribal nations over time

```
# Filter data for tribal requests with 'Denial' in requestResult
tribal_denials <- subset(data, tribalRequest == 1 & requestResult == "Denial")

# Count denials by year
denials_by_year <- table(tribal_denials$declarationYear)
denials_df <- as.data.frame(denials_by_year)
colnames(denials_df) <- c("Year", "Count")

# Convert Year to a numeric type for plotting
denials_df$Year <- as.numeric(as.character(denials_df$Year))

# Plot the data
ggplot(denials_df, aes(x = Year, y = Count)) +
  geom_line(color = "red") +
  geom_point(color = "red") +
  labs(title = "FEMA Aid Denials to Tribal Nations Over Time",
       x = "Year",
       y = "Number of Denials") +
  theme_minimal()
```

FEMA Aid Denials to Tribal Nations Over Time




Figure 4.2. Visualizing the number of FEMA aid acceptances over time to tribal nations shows a general upward trajectory, with a sharp spike in 2020, likely due to the COVID-19 pandemic. There appear to be far fewer denials than acceptances.

Examine FEMA aid acceptances to tribal nations over time, segmented by program type declared

```
# Load additional necessary libraries
library(reshape2)
```

```
##
## Attaching package: 'reshape2'
```

```
## The following object is masked from 'package:tidyr':
##
## smiths
```

```

# Filter data for tribal requests only
tribal_data <- subset(data, tribalRequest == 1)

# Aggregate the data by year and program acceptance
program_acceptances <- aggregate(cbind(ihProgramDeclared, iaProgramDeclared, paProgramDe
clared, hmProgramDeclared) ~ declarationYear,
                                data = tribal_data,
                                FUN = sum)

# Melt the data for plotting (long format)
melted_data <- melt(program_acceptances, id.vars = "declarationYear",
                    variable.name = "ProgramType",
                    value.name = "Acceptances")

# Map the column names to more descriptive labels
melted_data$ProgramType <- factor(melted_data$ProgramType,
                                  levels = c("ihProgramDeclared", "iaProgramDeclared",
"paProgramDeclared", "hmProgramDeclared"),
                                  labels = c("Individuals and Households Program Declare
d",
                                             "Individual Assistance Program Declared",
                                             "Public Assistance Program Declared",
                                             "Hazard Mitigation Program Declared"))

# Plot the data
ggplot(melted_data, aes(x = declarationYear, y = Acceptances, color = ProgramType)) +
  geom_line() +
  geom_point() +
  labs(title = "FEMA Aid Acceptances to Tribal Nations Over Time by Program Type",
        x = "Year",
        y = "Number of Acceptances",
        color = "Program Type") +
  theme_minimal()

```

```

## `geom_line()`: Each group consists of only one observation.
## i Do you need to adjust the group aesthetic?

```





Figure 5.1

Examine FEMA aid denials to tribal nations over time, segmented by program type requested

```

# Filter the data for tribal requests and denials
tribal_denials <- subset(data, tribalRequest == 1 & requestResult == "Denial")

# Aggregate the number of denials by year and program requested
program_denials <- aggregate(cbind(ihProgramRequested, iaProgramRequested, paProgramRequested, hmProgramRequested) ~ declarationYear,
                             data = tribal_denials,
                             FUN = sum)

# Melt the data for plotting (long format)
melted_denials <- melt(program_denials, id.vars = "declarationYear",
                      variable.name = "ProgramType",
                      value.name = "Denials")

# Map the column names to more descriptive labels
melted_denials$ProgramType <- factor(melted_denials$ProgramType,
                                     levels = c("ihProgramRequested", "iaProgramRequested", "paProgramRequested", "hmProgramRequested"),
                                     labels = c("Individuals and Households Program Requested",
                                                "Individual Assistance Program Requested",
                                                "Public Assistance Program Requested",
                                                "Hazard Mitigation Program Requested"))

# Plot the data
ggplot(melted_denials, aes(x = declarationYear, y = Denials, color = ProgramType)) +
  geom_line() +
  geom_point() +
  labs(title = "FEMA Aid Denials to Tribal Nations Over Time by Program Type",
       x = "Year",
       y = "Number of Denials",
       color = "Program Type") +
  theme_minimal()

```

```

## `geom_line()`: Each group consists of only one observation.
## i Do you need to adjust the group aesthetic?

```




Figure 5.2. Examining the number of acceptances and denials to tribal nations over time, segmented by program type, allows for a broader understanding of what types of aid are being requested.

Examining denials to tribal nations by state

```
# Filter data for tribal requests with 'Denial' in requestResult
tribal_denials <- subset(data, tribalRequest == 1 & requestResult == "Denial")

# Count denials by state
denials_by_state <- table(tribal_denials$state)
denials_df <- as.data.frame(denials_by_state)
colnames(denials_df) <- c("State", "Count")

# Sort the data frame by Count in descending order
denials_df <- denials_df[order(-denials_df$Count), ]

# Plot the data using ggplot2
ggplot(denials_df, aes(x = reorder(State, -Count), y = Count, fill = State)) +
  geom_bar(stat = "identity") +
  labs(title = "FEMA Aid Denials to Tribal Nations by State",
       x = "State",
       y = "Number of Denials") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```




Figure 6.1

Examining acceptances to tribal nation by state

```
# Filter data for tribal requests with 'Acceptance' in requestResult
tribal_acceptances <- subset(data, tribalRequest == 1 & requestResult == "Acceptance")

# Count acceptances by state
acceptances_by_state <- table(tribal_acceptances$state)
acceptances_df <- as.data.frame(acceptances_by_state)
colnames(acceptances_df) <- c("State", "Count")

# Sort the data frame by Count in descending order
acceptances_df <- acceptances_df[order(-acceptances_df$Count), ]

# Plot the data using ggplot2
ggplot(acceptances_df, aes(x = reorder(State, -Count), y = Count, fill = State)) +
  geom_bar(stat = "identity") +
  labs(title = "FEMA Aid Acceptances to Tribal Nations by State",
       x = "State",
       y = "Number of Acceptances") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

FEMA Aid Acceptances to Tribal Nations by State




Figure 6.2. Examining the number of acceptances and denials to tribal nations over time, segmented by state, allows for a broader overview of where aid is being requested. Florida leads aid acceptances to tribal nations, and South Dakota leads aid denials.

Examine acceptances to tribal nations by incident type

```
# Filter data for tribal requests with 'Acceptance' in requestResult
tribal_acceptances <- subset(data, tribalRequest == 1 & requestResult == "Acceptance")

# Count acceptances by incident type
acceptances_by_incident <- table(tribal_acceptances$incidentType)
acceptances_df <- as.data.frame(acceptances_by_incident)
colnames(acceptances_df) <- c("IncidentType", "Count")

# Sort the data frame by Count in descending order
acceptances_df <- acceptances_df[order(-acceptances_df$Count), ]

# Plot the data using ggplot2
ggplot(acceptances_df, aes(x = reorder(IncidentType, -Count), y = Count, fill = Incident
Type)) +
  geom_bar(stat = "identity") +
  labs(title = "FEMA Aid Acceptances to Tribal Nations by Incident Type",
       x = "Incident Type",
       y = "Number of Acceptances") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```




Figure 7.1

Examining denials to tribal nations by incident type

```
# Filter data for tribal requests with 'Denial' in requestResult
tribal_denials <- subset(data, tribalRequest == 1 & requestResult == "Denial")

# Count denials by incident type
denials_by_incident <- table(tribal_denials$incidentType)
denials_df <- as.data.frame(denials_by_incident)
colnames(denials_df) <- c("IncidentType", "Count")

# Sort the data frame by Count in descending order
denials_df <- denials_df[order(-denials_df$Count), ]

# Plot the data using ggplot2
ggplot(denials_df, aes(x = reorder(IncidentType, -Count), y = Count, fill = IncidentType)) +
  geom_bar(stat = "identity") +
  labs(title = "FEMA Aid Denials to Tribal Nations by Incident Type",
       x = "Incident Type",
       y = "Number of Denials") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

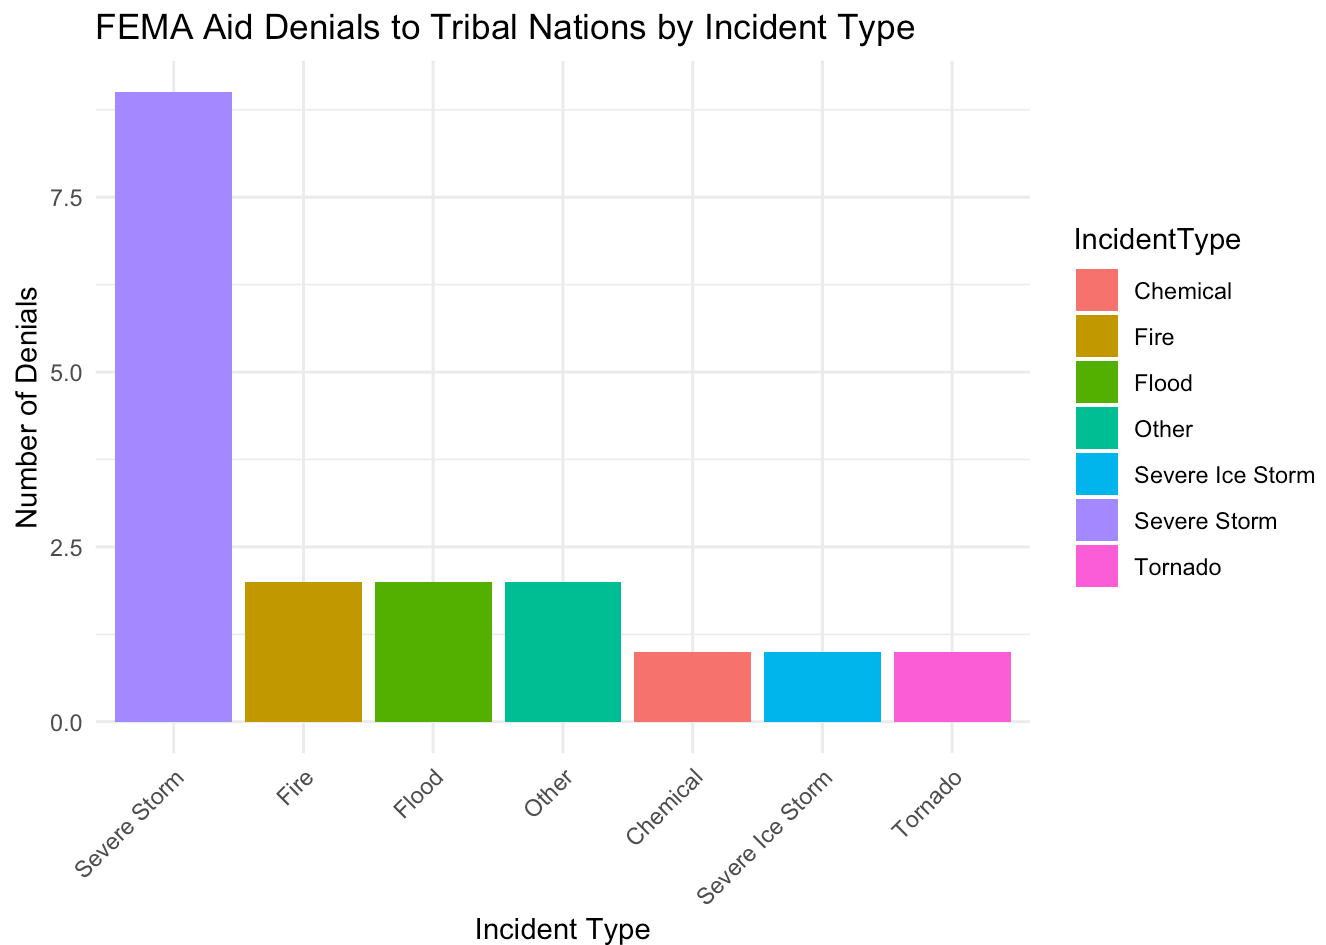


Figure 7.2. Examining the number of acceptances and denials to tribal nations over time, segmented by incident type, allows for a broader overview of why aid is being requested. Biological incidents are most frequently rewarded aid, and Severe Storm incidents are most frequently denied aid.

Examine the rate of denials per incident type, segmented by tribal nation vs. non-tribal nation, and scale results based on the total number of requests

```
# Calculate total requests by incident type and tribal status
total_requests <- data %>%
  group_by(incidentType, tribalRequest) %>%
  summarise(TotalRequests = n())
```

```
## `summarise()` has grouped output by 'incidentType'. You can override using the
## `.groups` argument.
```

```
# Filter data for denials
denials_data <- subset(data, requestResult == "Denial")

# Calculate number of denials by incident type and tribal status
denials_count <- denials_data %>%
  group_by(incidentType, tribalRequest) %>%
  summarise(Denials = n())
```

```
## `summarise()` has grouped output by 'incidentType'. You can override using the  
## `.groups` argument.
```

```
# Merge total requests and denials data  
merged_data <- merge(total_requests, denials_count,  
                     by = c("incidentType", "tribalRequest"),  
                     all.x = TRUE)  
  
# Replace NA values with 0 (for cases with no denials)  
merged_data$Denials[is.na(merged_data$Denials)] <- 0  
  
# Calculate the denial rate  
merged_data$DenialRate <- merged_data$Denials / merged_data$TotalRequests  
  
# Map tribalRequest values to labels  
merged_data$TribalStatus <- factor(merged_data$tribalRequest,  
                                   levels = c(0, 1),  
                                   labels = c("Non-Tribal", "Tribal"))  
  
# Plot the data using ggplot2  
ggplot(merged_data, aes(x = incidentType, y = DenialRate, fill = TribalStatus)) +  
  geom_bar(stat = "identity", position = "dodge") +  
  labs(title = "Scaled Denial Rate by Incident Type for Tribal vs. Non-Tribal Requests",  
        x = "Incident Type",  
        y = "Denial Rate",  
        fill = "Tribal Status") +  
  theme_minimal() +  
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

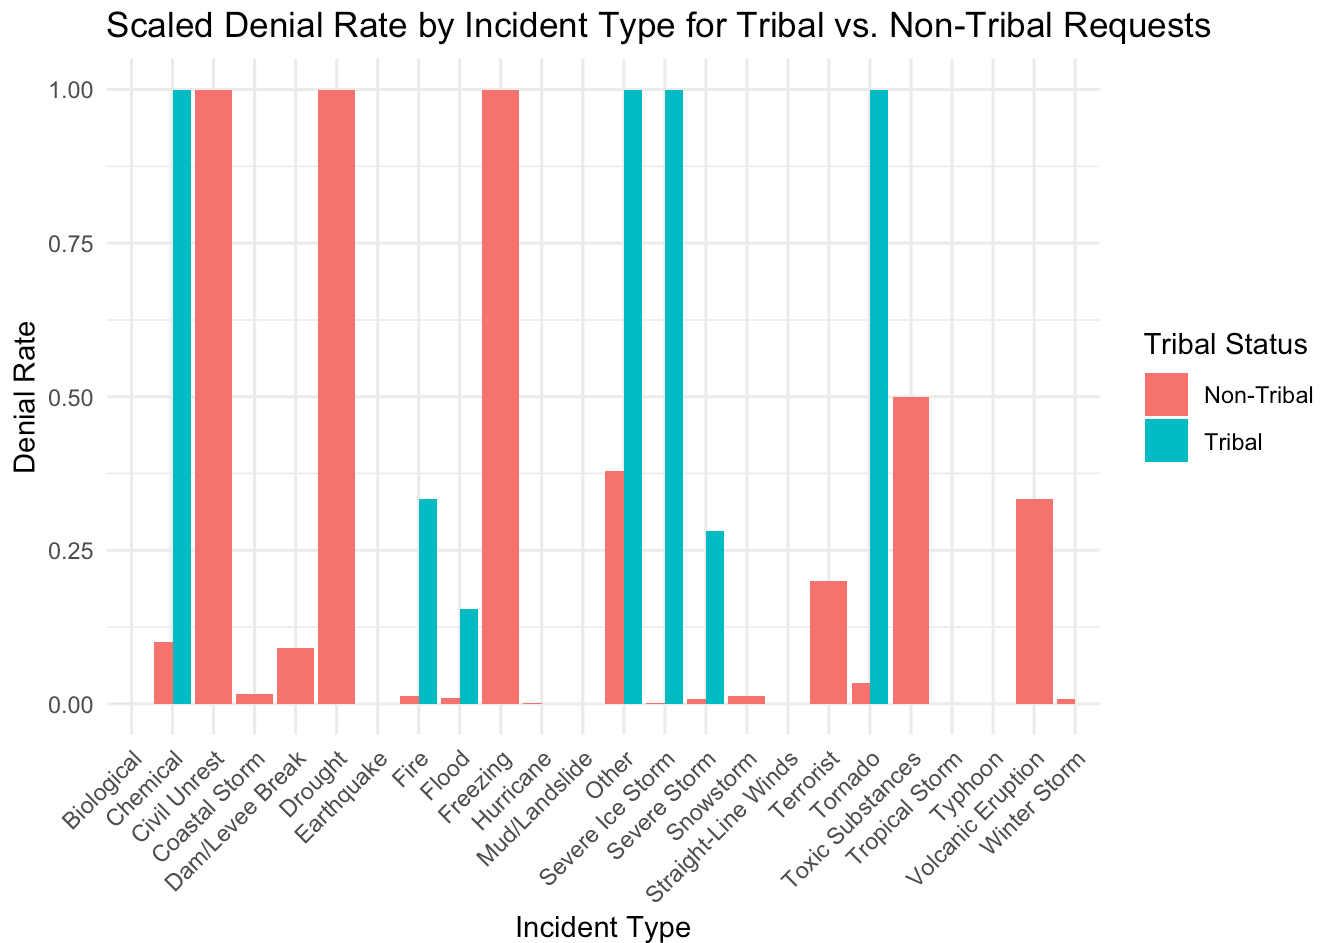



Figure 8.1

Examine the rate of acceptances per incident type, segmented by tribal nation vs. non-tribal nation, and scale results based on the total number of requests

```
# Calculate total requests by incident type and tribal status
total_requests <- data %>%
  group_by(incidentType, tribalRequest) %>%
  summarise(TotalRequests = n())
```

```
## `summarise()` has grouped output by 'incidentType'. You can override using the
## `.groups` argument.
```

```
# Filter data for acceptances
acceptances_data <- subset(data, requestResult == "Acceptance")

# Calculate number of acceptances by incident type and tribal status
acceptances_count <- acceptances_data %>%
  group_by(incidentType, tribalRequest) %>%
  summarise(Acceptances = n())
```

```
## `summarise()` has grouped output by 'incidentType'. You can override using the
## `.groups` argument.
```

```
# Merge total requests and acceptances data
merged_data <- merge(total_requests, acceptances_count,
                     by = c("incidentType", "tribalRequest"),
                     all.x = TRUE)

# Replace NA values with 0 (for cases with no acceptances)
merged_data$Acceptances[is.na(merged_data$Acceptances)] <- 0

# Calculate the acceptance rate
merged_data$AcceptanceRate <- merged_data$Acceptances / merged_data$TotalRequests

# Map tribalRequest values to labels
merged_data$TribalStatus <- factor(merged_data$tribalRequest,
                                   levels = c(0, 1),
                                   labels = c("Non-Tribal", "Tribal"))

# Plot the data using ggplot2
ggplot(merged_data, aes(x = incidentType, y = AcceptanceRate, fill = TribalStatus)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(title = "Scaled Acceptance Rate by Incident Type for Tribal vs. Non-Tribal Requests",
       x = "Incident Type",
       y = "Acceptance Rate",
       fill = "Tribal Status") +
  theme_minimal() +
  theme(axis.text.x = element_text(angle = 45, hjust = 1))
```

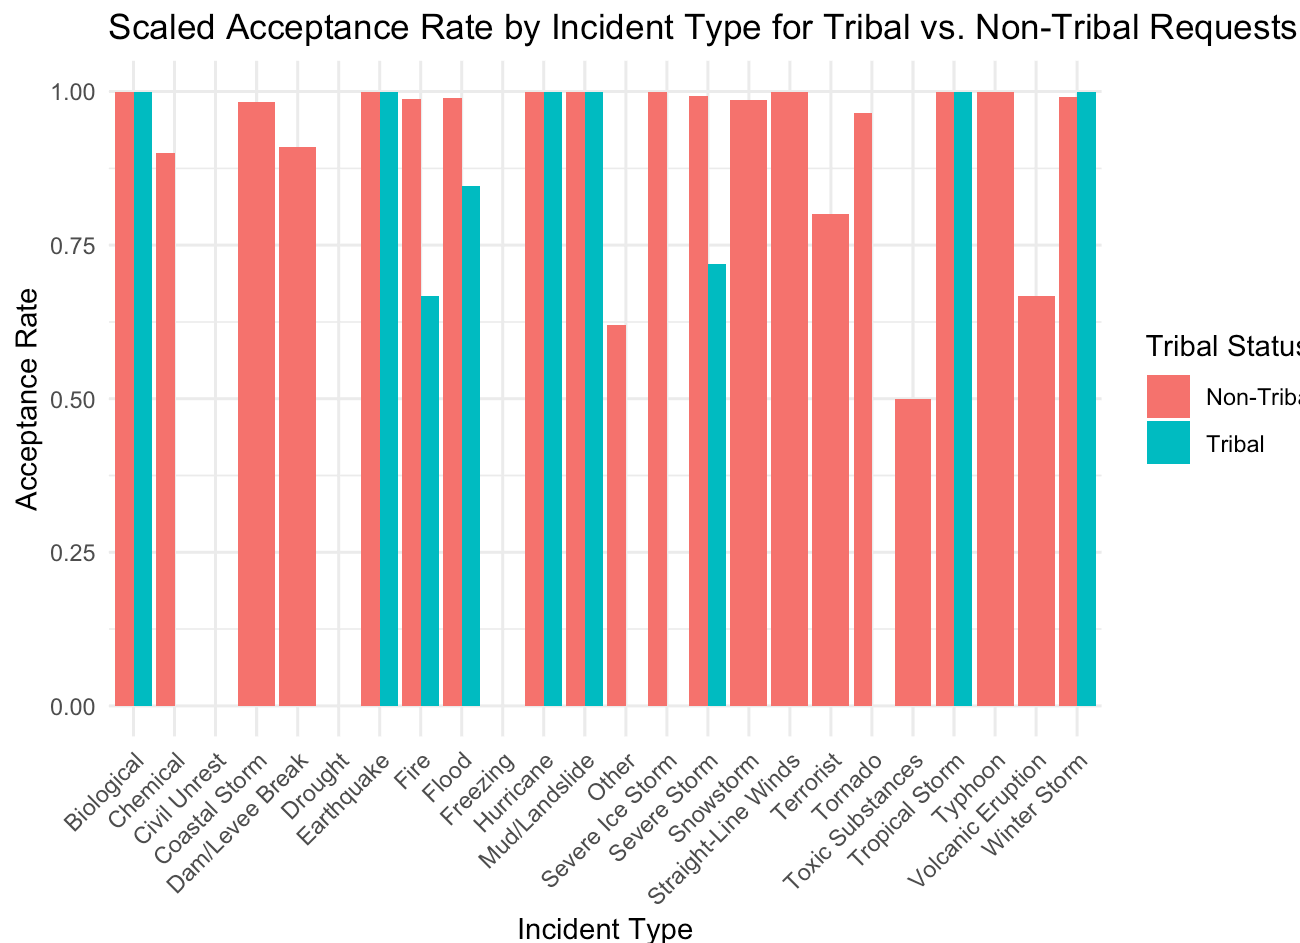


Figure 8.2. Examining the scaled denial rate by incident type for tribal vs. non-tribal requests, tribal requests are denied at a higher rate than non-tribal requests for all incident types for which both types of requests have been made.

That concludes the exploratory data analysis. Now, we will begin building a model.

BINARY LOGISTIC REGRESSION

Our goal is to analyze the likelihood that a FEMA aid request will be denied. In particular, we want to see if the coefficient for whether it was a tribal request is high and statistically significant.

First, some preliminary steps:

```
# Import additional relevant libraries
library(caret)
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':
##
## cov, smooth, var
```

```
library(dplyr)
```

```
# Clean data
# make dependent variable (request result) binary - 1 for accepted, 0 for denied
# convert into factor for classification
data$requestResult_binary <- ifelse(data$requestResult == "Acceptance",1,0)
data$requestResult_binary <- factor(data$requestResult_binary, levels = c(0, 1))
```

```
# Convert incident type to numeric by manually assigning numbers to categories
data$incidentType_numeric <- as.numeric(factor(data$incidentType,
                                              levels = c(unique(data$incidentType))))

# Convert state to numeric by manually assigning numbers to categories
data$state_numeric <- as.numeric(factor(data$state,
                                       levels = c(unique(data$state))))
```

```
# only keep potentially relevant columns
data_clean <- select(data, tribalRequest, incidentType_numeric, state_numeric, region, requestResult_binary)
head(data_clean)
```

	tribalRequest <int>	incidentType_numeric <dbl>	state_numeric <dbl>	region <int>	requestResult_binary <fct>
152	0	1	1	10	1
153	0	1	1	10	1
154	0	1	1	10	1
155	0	1	1	10	1
156	0	1	1	10	1
157	0	1	1	10	1
6 rows					

Then, balance our data and split into train/test sets.

```
table(data_clean$requestResult_binary)
```

```
##
##      0      1
## 136 24303
```

```
#note: there are 136 denials, 24226 acceptances
```

```
# filter for 0s and 1s
zeros <- data_clean[data_clean$requestResult_binary == 0,]
ones <- data_clean[data_clean$requestResult_binary == 1,]

# sample equal number of 0's and 1's
set.seed(35) # for reproducibility
sample_zeros <- zeros[sample(nrow(zeros), 136), ]
sample_ones <- ones[sample(nrow(ones), 136),]

# merge the sample datasets
data_balanced <- rbind(sample_zeros, sample_ones)

# confirm done correctly
table(data_balanced$requestResult_binary) # yes! :)
```

```
##
##    0    1
## 136 136
```

Now, split data into training/testing data.

```
set.seed(57) # for reproducibility
splitIndex <- createDataPartition(data_balanced$requestResult_binary, p = 0.7, list = FALSE)

train_data <- data_balanced[splitIndex, ]
test_data <- data_balanced[-splitIndex, ]
```

Now, we will do backward variable selection to only keep the most relevant variables in the binary logistic regression model that we train.

```
# initialize model w/ all predictors included
full_model <- glm(requestResult_binary ~ ., data = train_data, family = binomial(link="logit"))

# view model with all predictors included
summary(full_model)
```

```
##
## Call:
## glm(formula = requestResult_binary ~ ., family = binomial(link = "logit"),
##      data = train_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.204704    0.564341  -0.363   0.7168
## tribalRequest  -2.642508    1.065964  -2.479   0.0132 *
## incidentType_numeric  0.024995    0.029465   0.848   0.3963
## state_numeric    0.017905    0.009738   1.839   0.0660 .
## region         -0.064056    0.070287  -0.911   0.3621
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 266.17  on 191  degrees of freedom
## Residual deviance: 246.06  on 187  degrees of freedom
## AIC: 256.06
##
## Number of Fisher Scoring iterations: 5
```

Figure 9.

```
# iteratively remove predictors until model no longer improves
backward_model <- step(full_model, direction = "backward")
```

```
## Start: AIC=256.06
## requestResult_binary ~ tribalRequest + incidentType_numeric +
##   state_numeric + region
##
##           Df Deviance   AIC
## - incidentType_numeric 1  246.79 254.79
## - region                1  246.90 254.90
## <none>                  246.06 256.06
## - state_numeric         1  249.50 257.50
## - tribalRequest         1  257.33 265.33
##
## Step: AIC=254.79
## requestResult_binary ~ tribalRequest + state_numeric + region
##
##           Df Deviance   AIC
## - region          1  247.63 253.63
## <none>             246.79 254.79
## - state_numeric   1  250.25 256.25
## - tribalRequest   1  258.63 264.63
##
## Step: AIC=253.63
## requestResult_binary ~ tribalRequest + state_numeric
##
##           Df Deviance   AIC
## <none>             247.63 253.63
## - state_numeric   1  251.77 255.77
## - tribalRequest   1  262.47 266.47
```

```
# view model
summary(backward_model)
```

```
##
## Call:
## glm(formula = requestResult_binary ~ tribalRequest + state_numeric,
##      family = binomial(link = "logit"), data = train_data)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -0.433344   0.325481  -1.331   0.18306
## tribalRequest -2.848205   1.051010  -2.710   0.00673 **
## state_numeric  0.019251   0.009576   2.010   0.04438 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 266.17  on 191  degrees of freedom
## Residual deviance: 247.63  on 189  degrees of freedom
## AIC: 253.63
##
## Number of Fisher Scoring iterations: 5
```

Figure 10

Next steps are testing our model to see how it performs.

```
# Note: The training data does not include all types of disasters that the testing set might, since they were split based on a randomized process. To address this, we will convert those oddball results into "Other" so that we can still use our testing data!
test_data_adj <- test_data

unique_test_values <- setdiff(unique(test_data_adj$incidentType_numeric), unique(train_data$incidentType_numeric))

test_data_adj$incidentType <- ifelse(test_data_adj$incidentType_numeric %in% unique_test_values, "Other", test_data_adj$incidentType_numeric)

predicted_probabilities <- predict(backward_model,
                                 newdata = test_data_adj,
                                 type = "response")
predicted_classes <- ifelse(predicted_probabilities > 0.5, 1, 0) #threshold = 0.5

cmatrix <- table(test_data_adj$requestResult_binary, predicted_classes)
print(cmatrix)
```

```
##      predicted_classes
##      0  1
## 0 22 18
## 1 17 23
```

Figure 11

MODEL TUNING AND EVALUATION

To better understand how our model's performance changes as we introduce more predictors, we will take a look at the Validation Curve, using LOOCV. This will allow us to identify if our model is overfitting or underfitting with various numbers of predictors. In addition, using LOOCV specifically will allow us to train our model more robustly, given that we have a small dataset to work with.

```
# LOOCV function:
#   traindata = training dataset
#   target = the column to predict
#   parameters = list of parameters to use in model
loocv <- function(traindata, target, parameters) {
  # for loocv, m is the size of the training data
  m = nrow(traindata)
  train_control = trainControl(method="cv", number=m)

  # train the model
  model <- train(as.formula(paste(target, parameters)),
                 data = traindata,
                 method = "glm",
                 family = binomial(link = "logit"),
                 trControl = train_control)

  return(model)
}
```

Note: When training, the code may output “Warning: There were missing values in resampled performance measures.” This is most likely due to a resample where one of our result classes has 0 samples, causing the code to be unable to calculate a performance metric that involves looking at both classes.

```
# Train the model using LOOCV. Create four iterations of the model with decreasing number of predictors.
cv_model_full <- loocv(train_data, "requestResult_binary", "~ tribalRequest + state_numeric + incidentType_numeric + region")
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

```
cv_model_three <- loocv(train_data, "requestResult_binary", "~ tribalRequest + state_numeric + region")
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

```
cv_model_two <- loocv(train_data, "requestResult_binary", "~ tribalRequest + state_numeric")
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

```
cv_model_one <- loocv(train_data, "requestResult_binary", "~ tribalRequest")
```

```
## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.
```

Next we will evaluate the performance of each model iteration, collecting the training and testing accuracies for each model along the way.

```
# Create a dataframe to hold the results to plot
validCurve <- data.frame(predictors = integer(4),
                          trainAccuracy = integer(4),
                          testAccuracy = integer(4))
```

Model equation: requestResult_binary ~ tribalRequest

```
# Evaluate the model with one predictor
validCurve$predictors[1] <- 1
validCurve$trainAccuracy[1] <- cv_model_one$results$Accuracy

# Predict on testing data
test_data_one <- test_data_adj %>% dplyr::select(requestResult_binary, tribalRequest)
predicted_probabilities <- predict(cv_model_one,
                                 newdata = test_data_one,
                                 type = "prob")
predicted_classes <- ifelse(predicted_probabilities$'1' > 0.5, 1, 0) #threshold = 0.5
predicted_classes <- factor(predicted_classes, levels=c(0,1))

correct_predictions <- sum(predicted_classes == test_data_one$requestResult_binary)
total_observations <- nrow(test_data_one)
validCurve$testAccuracy[1] <- correct_predictions / total_observations
```

Model equation: requestResult_binary ~ tribalRequest + state_numeric

```
# Evaluate the model with two predictors
validCurve$predictors[2] <- 2
validCurve$trainAccuracy[2] <- cv_model_two$results$Accuracy

# Predict on testing data
test_data_two <- test_data_adj %>% dplyr::select(requestResult_binary, tribalRequest, state_numeric)
predicted_probabilities <- predict(cv_model_two,
                                 newdata = test_data_two,
                                 type = "prob")
predicted_classes <- ifelse(predicted_probabilities$'1' > 0.5, 1, 0) #threshold = 0.5
predicted_classes <- factor(predicted_classes, levels=c(0,1))

correct_predictions <- sum(predicted_classes == test_data_two$requestResult_binary)
total_observations <- nrow(test_data_two)
validCurve$testAccuracy[2] <- correct_predictions / total_observations
```

Model equation: requestResult_binary ~ tribalRequest + state_numeric + region

```
# Evaluate the model with three predictors
validCurve$predictors[3] <- 3
validCurve$trainAccuracy[3] <- cv_model_three$results$Accuracy

# Predict on testing data
test_data_three <- test_data_adj %>% dplyr::select(requestResult_binary, tribalRequest,
state_numeric, region)
predicted_probabilities <- predict(cv_model_three,
                                newdata = test_data_three,
                                type = "prob")
predicted_classes <- ifelse(predicted_probabilities$'1' > 0.5, 1, 0) #threshold = 0.5
predicted_classes <- factor(predicted_classes, levels=c(0,1))

correct_predictions <- sum(predicted_classes == test_data_three$requestResult_binary)
total_observations <- nrow(test_data_three)
validCurve$testAccuracy[3] <- correct_predictions / total_observations
```

Model equation: requestResult_binary ~ tribalRequest + state_numeric + incidentType_numeric + region

```
# Evaluate the model with all predictors
validCurve$predictors[4] <- 4
validCurve$trainAccuracy[4] <- cv_model_full$results$Accuracy

# Predict on testing data
test_data_full <- test_data_adj %>% dplyr::select(requestResult_binary, tribalRequest, s
tate_numeric, region, incidentType_numeric)
predicted_probabilities <- predict(cv_model_full,
                                newdata = test_data_full,
                                type = "prob")
predicted_classes <- ifelse(predicted_probabilities$'1' > 0.5, 1, 0) #threshold = 0.5
predicted_classes <- factor(predicted_classes, levels=c(0,1))

correct_predictions <- sum(predicted_classes == test_data_full$requestResult_binary)
total_observations <- nrow(test_data_full)
validCurve$testAccuracy[4] <- correct_predictions / total_observations
```

```
# View the results
head(validCurve)
```

	predictors <dbl>	trainAccuracy <dbl>	testAccuracy <dbl>
1	1	0.5677083	0.5500
2	2	0.6145833	0.5625
3	3	0.5885417	0.5625
4	4	0.6041667	0.5875

4 rows

```
# Plot the validation curve
ggplot(validCurve, aes(x = predictors)) +
  geom_line(aes(y = trainAccuracy, col = "Training Accuracy")) +
  geom_point(aes(y = trainAccuracy, col = "Training Accuracy")) +
  geom_line(aes(y = testAccuracy, col = 'Testing Accuracy')) +
  geom_point(aes(y = testAccuracy, col='Testing Accuracy')) +
  labs(title = "Validation Curve for Logistic Regression of FEMA Acceptance or Denial",
       x = "Model Parameters",
       y = "Accuracy") +
  theme_minimal() +
  ylim(0,1) +
  guides(fill = guide_legend(reverse=TRUE)) +
  scale_color_discrete(name = "")
```

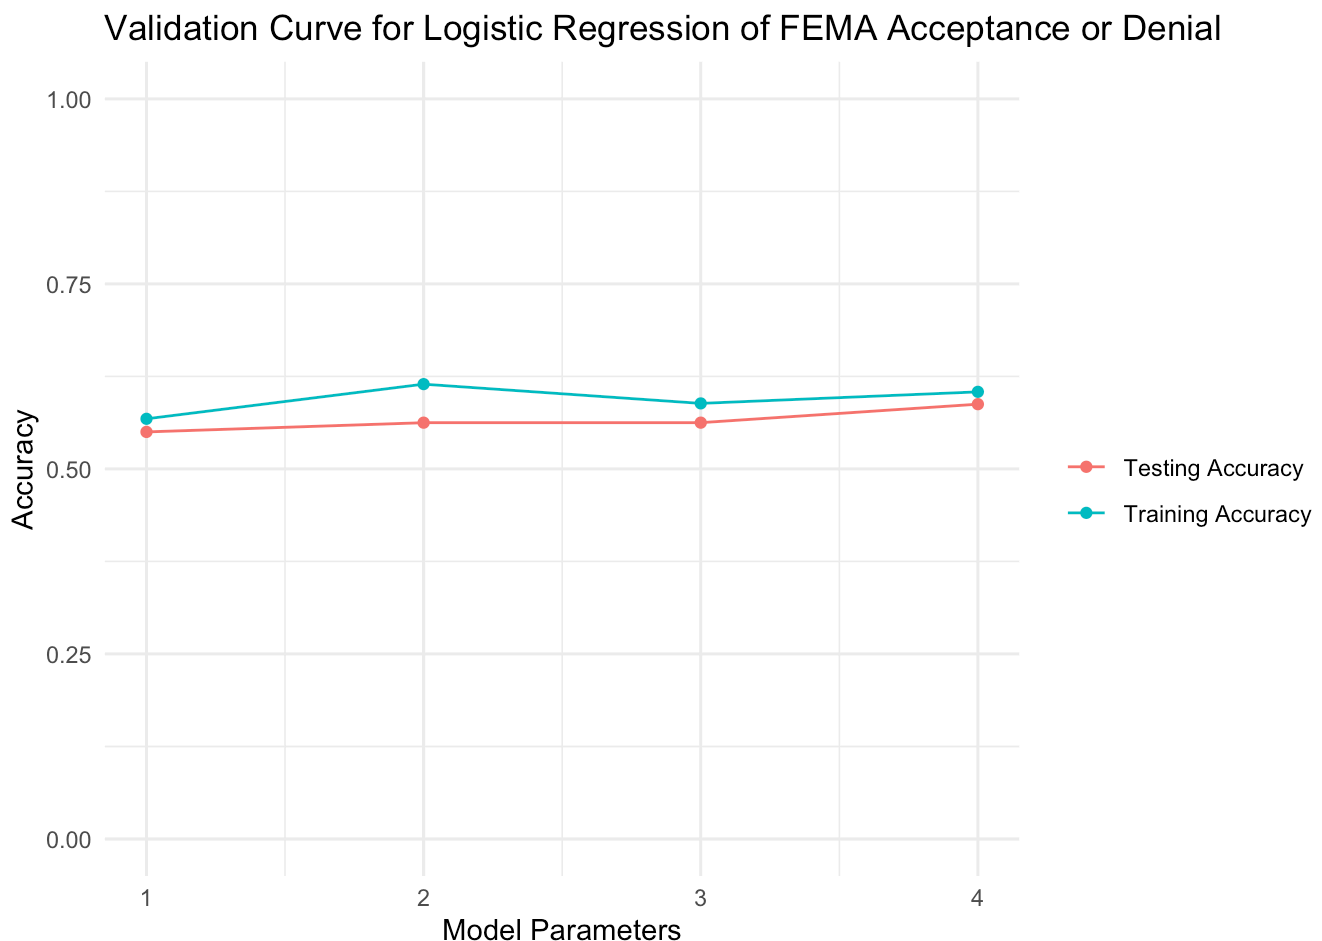


Figure 12

As we can see in the validation curve, there is not significant variation in the testing vs. training accuracies for the differing number of model parameters. We can see that the testing accuracy and training accuracy are closest in value for the model with 1 parameter and the model with 4 parameters. We will continue analyzing the model with 4 parameters to get a better sense of its overall performance.

For our final model, we will take a look at some classification metrics to evaluate its overall performance. We will calculate the metrics based on the correct classification of the accepted request class.

Here we see that the accuracy is 0.56, with 45 samples classified correctly and 35 samples classified incorrectly.

```
# Confusion Matrix
cm <- confusionMatrix(predicted_classes, test_data_full$requestResult_binary)
cm
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 21 14
##           1 19 26
##
##           Accuracy : 0.5875
##           95% CI : (0.4718, 0.6965)
##       No Information Rate : 0.5
##       P-Value [Acc > NIR] : 0.07282
##
##           Kappa : 0.175
##
##  Mcnemar's Test P-Value : 0.48623
##
##           Sensitivity : 0.5250
##           Specificity : 0.6500
##           Pos Pred Value : 0.6000
##           Neg Pred Value : 0.5778
##           Prevalence : 0.5000
##           Detection Rate : 0.2625
##       Detection Prevalence : 0.4375
##           Balanced Accuracy : 0.5875
##
##           'Positive' Class : 0
##
```

Figure 13

The precision is 0.56, meaning that 56% of the samples that the model classified as accepted requests were correctly identified as accepted requests. The recall is 0.55, meaning that 55% of the samples that are actually accepted requests were correctly identified as accepted requests. The F1 score, which takes into account both precision and recall, is 0.56.

```
# Overall evaluation metrics
cm$byClass
```

##	Sensitivity	Specificity	Pos Pred Value
##	0.5250000	0.6500000	0.6000000
##	Neg Pred Value	Precision	Recall
##	0.5777778	0.6000000	0.5250000
##	F1	Prevalence	Detection Rate
##	0.5600000	0.5000000	0.2625000
##	Detection Prevalence	Balanced Accuracy	
##	0.4375000	0.5875000	

Figure 14

Finally, we can see in the model summary that the tribalRequest parameter is statistically significant with a p-value of 0.0132.

```
# Look at model summary
summary(cv_model_full)
```

```
##
## Call:
## NULL
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.204704   0.564341  -0.363   0.7168
## tribalRequest  -2.642508   1.065964  -2.479   0.0132 *
## state_numeric    0.017905   0.009738   1.839   0.0660 .
## incidentType_numeric 0.024995   0.029465   0.848   0.3963
## region         -0.064056   0.070287  -0.911   0.3621
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##    Null deviance: 266.17  on 191  degrees of freedom
## Residual deviance: 246.06  on 187  degrees of freedom
## AIC: 256.06
##
## Number of Fisher Scoring iterations: 5
```

Figure 15

OVERALL LOGISTIC REGRESSION CONCLUSIONS

Overall, the metrics above mean that our model is not currently able to learn how to classify FEMA requests into acceptances or denials. While we tried to combat this with our use of LOOCV, this is most likely due to our dataset being too small to adequately train our model with. However, the finding that the tribalRequest parameter is statistically significant is important to our research question in that it indicates that whether or not a request was made by a Tribal Nation is significant to if the request is denied or accepted.

RANDOM FOREST MODEL

Given that our accuracy is fairly low from a binary logistic regression model, we will employ a random forest model to attempt to better learn how to classify FEMA requests into acceptances or denials.

```
# Load necessary libraries
library(tidyverse)
library(randomForest)
```

```
## randomForest 4.7-1.2
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:dplyr':
##
##      combine
```

```
## The following object is masked from 'package:ggplot2':
##
##      margin
```

```
library(caret)
library(ggplot2)
```

```
# Create new columns for the model
data <- data %>%
  mutate(
    requestResult_binary = factor(ifelse(requestResult == "Acceptance", 1, 0), levels =
c(0, 1)),
    tribalRequest_factor = as.factor(tribalRequest),
    incidentType_numeric = as.numeric(factor(incidentType)),
    state_numeric = as.numeric(factor(state))
  )

# Select relevant columns for the model
data_RF <- select(data, tribalRequest_factor, incidentType_numeric, state_numeric, region, requestResult_binary)

# View first few rows of Random Forest dataset
head(data_RF)
```

tribalRequest_factor <fct>	incidentType_numeric <dbl>	state_numeric <dbl>	region <int>	requestResult_binary <fct>
152 0	9	1	10	1
153 0	9	1	10	1
154 0	9	1	10	1
155 0	9	1	10	1
156 0	9	1	10	1
157 0	9	1	10	1
6 rows				

```
# Balance the dataset
# filter for 0s and 1s
zeros <- data_RF[data_RF$requestResult_binary == 0,]
ones <- data_RF[data_RF$requestResult_binary == 1,]

# sample equal number of 0's and 1's
set.seed(35) # for reproducibility
sample_zeros <- zeros[sample(nrow(zeros), 136), ]
sample_ones <- ones[sample(nrow(ones), 136),]

# merge the sample datasets
data_RF_balanced <- rbind(sample_zeros, sample_ones)

# confirm done correctly
table(data_RF_balanced$requestResult_binary) # yes! :)
```

```
##
##    0    1
## 136 136
```

```
# Split the data into training and test sets
set.seed(123)
train_index <- createDataPartition(data_RF_balanced$requestResult_binary, p = 0.7, list
= FALSE)
train_data <- data_RF_balanced[train_index, ]
test_data <- data_RF_balanced[-train_index, ]
```

```
# Train the Random Forest model
rf_model <- randomForest(requestResult_binary ~ ., data = train_data, importance = TRUE,
ntree = 500)

# Model summary
print(rf_model)
```



```
##
## Call:
##  randomForest(formula = requestResult_binary ~ ., data = train_data,      importance
= TRUE, ntree = 500)
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 2
##
##          OOB estimate of  error rate: 22.4%
## Confusion matrix:
##      0  1 class.error
## 0 76 20   0.2083333
## 1 23 73   0.2395833
```

Figure 16

```
# Predict on the test data
predictions <- predict(rf_model, test_data)

# Confusion Matrix
confusion <- confusionMatrix(predictions, test_data$requestResult_binary)
print(confusion)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0  1
##           0 31 15
##           1  9 25
##
##              Accuracy : 0.7
##              95% CI : (0.5872, 0.7974)
##      No Information Rate : 0.5
##      P-Value [Acc > NIR] : 0.0002258
##
##              Kappa : 0.4
##
##  Mcnemar's Test P-Value : 0.3074342
##
##              Sensitivity : 0.7750
##              Specificity : 0.6250
##              Pos Pred Value : 0.6739
##              Neg Pred Value : 0.7353
##              Prevalence : 0.5000
##              Detection Rate : 0.3875
##      Detection Prevalence : 0.5750
##              Balanced Accuracy : 0.7000
##
##              'Positive' Class : 0
##
```

Figure 17

```
# Accuracy
accuracy <- confusion$overall['Accuracy']
cat("Accuracy:", accuracy, "\n")
```

```
## Accuracy: 0.7
```

```
# Feature Importance Visualization
importance <- as.data.frame(importance(rf_model))
importance <- importance %>% rownames_to_column(var = "Feature")
ggplot(importance, aes(x = reorder(Feature, MeanDecreaseGini), y = MeanDecreaseGini)) +
  geom_bar(stat = "identity", fill = "blue") +
  coord_flip() +
  labs(title = "Feature Importance (Mean Decrease Gini)", x = "Features", y = "Importance") +
  theme_minimal()
```




Figure 18

```
# Calculate Mean Squared Error (MSE)
mse <- mean((as.numeric(as.character(predictions)) - as.numeric(as.character(test_data$requestResult_binary)))^2)
cat("Mean Squared Error (MSE):", mse, "\n")
```

```
## Mean Squared Error (MSE): 0.3
```

```
# Plotting Error Rate Across Trees
```

```
error_rates <- data.frame(Trees = 1:rf_model$ntree,  
                          OOB_Error = rf_model$err.rate[, 1]) # OOB Error (1st column in  
err.rate)
```

```
ggplot(error_rates, aes(x = Trees, y = OOB_Error)) +  
  geom_line(color = "blue") +  
  labs(title = "OOB Error Rate Across Trees", x = "Number of Trees", y = "OOB Error Rat  
e") +  
  theme_minimal()
```

OOB Error Rate Across Trees




Figure 19. The optimal number of trees to reduce OOB error is around 100 trees.

```
# Train the Random Forest model with 125 trees
```

```
tuned_rf <- randomForest(requestResult_binary ~ ., data = train_data, importance = TRUE,  
ntree = 125)
```

```
# Model summary
```

```
print(tuned_rf)
```

```
##
## Call:
##  randomForest(formula = requestResult_binary ~ ., data = train_data,      importance
= TRUE, ntree = 125)
##              Type of random forest: classification
##              Number of trees: 125
## No. of variables tried at each split: 2
##
##          OOB estimate of  error rate: 23.44%
## Confusion matrix:
##      0  1 class.error
## 0 76 20    0.2083333
## 1 25 71    0.2604167
```

Figure 20

```
# Predict on the test data
predictions <- predict(tuned_rf, test_data)

# Confusion Matrix
confusion <- confusionMatrix(predictions, test_data$requestResult_binary)
print(confusion)
```

```
## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0  1
##           0 32 15
##           1  8 25
##
##              Accuracy : 0.7125
##              95% CI : (0.6005, 0.8082)
##   No Information Rate : 0.5
##   P-Value [Acc > NIR] : 9.156e-05
##
##              Kappa : 0.425
##
##  Mcnemar's Test P-Value : 0.2109
##
##              Sensitivity : 0.8000
##              Specificity : 0.6250
##              Pos Pred Value : 0.6809
##              Neg Pred Value : 0.7576
##              Prevalence : 0.5000
##              Detection Rate : 0.4000
##              Detection Prevalence : 0.5875
##              Balanced Accuracy : 0.7125
##
##              'Positive' Class : 0
##
```

Figure 21

```
# Accuracy
accuracy <- confusion$overall['Accuracy']
cat("Tuned Model Accuracy:", accuracy, "\n")
```

```
## Tuned Model Accuracy: 0.7125
```

```
# Feature Importance Visualization
importance <- as.data.frame(importance(tuned_rf))
importance <- importance %>% rownames_to_column(var = "Feature")
ggplot(importance, aes(x = reorder(Feature, MeanDecreaseGini), y = MeanDecreaseGini)) +
  geom_bar(stat = "identity", fill = "blue") +
  coord_flip() +
  labs(title = "Feature Importance (Mean Decrease Gini) - Tuned Model", x = "Features",
y = "Importance") +
  theme_minimal()
```




Figure 22

OVERALL RANDOM FOREST CONCLUSIONS

Overall, this model has a better accuracy than the previous logistic regression model, with a score of 0.7. The incident type was found to be the most important feature to split on within the random forest decision trees, followed by the state, region, and then if the request was a tribal request. The importance of the incident type compared to if the request was a tribal request is roughly 45 compared to 5, respectively. This result indicates that if the request was from a tribal nation was not found to be as significant of a determining factor in an acceptance or denial compared to the other features included.

The out-of-bag error rate graph shows that the error decreases from 0 to 100 trees, then increases around 200 to 350 trees, then decreases again for 400 to 500 trees. This indicates that the optimal number of trees to limit the out-of-bag error rate is around 100-125 or 500. To limit redundancy and increase efficiency, it would be best to use the smaller number of trees.

The model was rebuilt with 125 trees, and achieved an accuracy of 0.7. This result shows that the reduction in trees did not negatively affect the correct classification of the testing data, allowing us to reduce unnecessary complexity. The feature importances found for this new model were very similar to the initial model's, showing consistency in the results.